

# Updated: A trip advisor based on LightRAG

Fanshi Meng (904055546), Junbo Zou (904145436)]

## 1 Project Overview

Trip Advisor builds a knowledge-graph-augmented Retrieval-Augmented Generation (RAG)Lewis et al. (2020) travel assistant for Florida businesses using the **LightRAG**Zhu et al. (2024) framework. The system ingests Yelp-style business descriptions, constructs a graph of entities and relations, embeds text and graph elements into a vector store, and answers user queries via an LLM using graph- and chunk-aware context.

### High-level Architecture

- **Ingestion & Normalization:** Raw business JSON is transformed into natural-language records to preserve semantics for LLM extraction.
- **Graph Construction (GraphML):** LLM-driven entity and relation extraction yields a knowledge graph persisted in `graph_chunk_entity_relation.graphml`.
- **Vector Indexing (Nano-VectorDB):** Entities, relations, and chunks are embedded (1024-d) for semantic retrieval.
- **Query Layer (LightRAG Local Mode):** Top- $k$  entity retrieval → relation fan-out → chunk selection; context is assembled and routed to the LLM for answer generation.

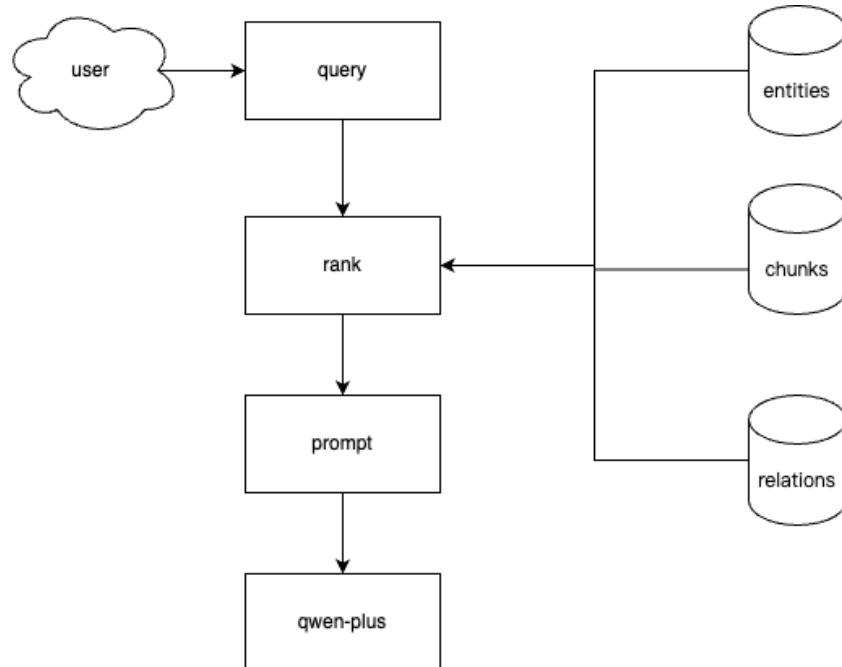


Figure 1: Trip Advisor overview

## 2 System Context & Existing Infrastructure

Our recommendation system based on the LightRAG technique integrates with an existing production web application<sup>1</sup> (deployed on AWS Amplify) backed by a PostgreSQL database that manages user travel plans and point-of-interest (POI) information.

### 2.1 Production Database Schema

The production system consists of five core tables, and our LightRAG system collects data from the following two tables:

**1. itineraries Table** Captures user travel plans with structured preference fields:

- **Destination:** destination\_city (e.g., “Tampa, FL”)
- **Travel Dates:** start\_date, end\_date
- **Budget:** budget\_in\_cents (e.g., 5000 = \$50 per activity)
- **Party Composition:** has\_children, has\_elderly, number\_of\_travelers
- **Travel Style:** activity\_intensity (LIGHT/ MODERATE/ INTENSE), travel\_pace (RELAXED/ MODERATE/ PACKED)
- **Preferences:** additional\_preferences (free text), seeded\_recommendations (JSONB)

**2. itinerary\_preferred\_categories Table** Links itineraries to user-selected categories (e.g., “Italian”, “Museums”, “Beach”) in a one-to-many relationship.

### 2.2 Current Workflow Without LightRAG

Without our knowledge-graph-augmented approach, the system relies purely on:

- **SQL Filtering:** Queries the places table by city, extracts metadata from JSONB, and applies basic WHERE clauses.
- **Limited Personalization:** No reasoning about *why* a POI suits a user’s itinerary preferences.

### 2.3 Integration Goal

Our LightRAG pipeline aims to *augment* this infrastructure by:

1. Constructing a semantic knowledge graph from places data to enable natural-language query understanding.
2. Leveraging itinerary preference fields (budget, family composition, activity intensity) to personalize recommendations.
3. Generating contextual, narrative-style explanations for *why* specific POIs match user needs.

This integration challenge—balancing structured database efficiency with LightRAG’s semantic reasoning—motivates the hybrid architecture discussed in Section 5.

---

<sup>1</sup><https://staging.d2ocft4tiyr0oy.amplifyapp.com/>

### 3 Progress

#### Knowledge Graph

The data source of Knowledge Graph is from Yelp. To be specific, we filter the data of Florida.

##### Entity Node Example

```
Entity: "Roman Forum"
Type: BUSINESS
Attributes:
- entity_name: "Roman Forum"
- entity_type: "BUSINESS"
- description: "American-Italian restaurant"
- source_id: "chunk-abc123"
- embedding: [0.12, -0.34, 0.56, ...]
```

##### Relation Edge Example

```
Relation: "Roman Forum" → "Tampa"
Type: LOCATED_IN
Attributes:
- src_id: "Roman Forum"
- tgt_id: "Tampa"
- keywords: "located, Tampa, Florida"
- description: "Restaurant is in Tampa"
- weight: 0.95
```

#### 3.1 Data Pipeline

1. **Preprocessing:** Convert each Yelp JSON record (e.g., business name, categories, hours, features) into a natural-language paragraph suitable for LLM extraction.
2. **Chunking:** Long concatenated text is windowed into  $\sim 17$  chunks of  $\sim 1200$ -1500 tokens each for stable extraction and retrieval.
3. **Entity/Relation Extraction:** The LLM (Qwen) extracts *BUSINESS*, *LOCATION*, *CATEGORY*, *FEATURE*, etc., and relations such as *LOCATED\_IN*, *HAS\_CATEGORY*, *OFFERS\_FEATURE*.
4. **Embedding + Storage:** All entities, relations, and chunks are embedded (1024-d) and written to Nano-VectorDB; the graph is saved as GraphML.

#### 3.2 Serving and Querying

We use LightRAG *local* mode for balanced quality and latency:

1. Embed the user query (e.g., "Recommend Italian restaurants in Tampa").
2. Retrieve top- $k$  similar entities (e.g., *Tampa*, *Italian Cuisine*, *Roman Forum*, *Joe's Pizza*); expand 1-2 hops over relations.

3. Select the most relevant chunks (up to 17) and construct a natural-language context.
4. Prompt the LLM with (system prompt + context + user query) to produce the final answer.

### 3.3 Engineering Work Completed

- Implemented batch ingestion (100 records/batch) and asynchronous inserts.
- Enabled LLM response caching to reduce repeated query cost and end-to-end latency.
- Wired Nano-VectorDB for cosine similarity search across entities/relations/chunks.
- Added detailed, inspectable retrieval traces for debugging (entities/relations/chunks and chunk IDs).

## 4 Preliminary Results

For the example query “Recommend Italian restaurants in Tampa”, the system retrieves relevant entities and produces grounded suggestions such as **Roman Forum** (American-Italian) and **Joe’s Pizza** (NY-style pizza), including ratings, addresses, operating hours (when available), and short justifications. In typical local-mode retrieval, the system returns approximately ~40 entities and ~80–90 relations before deduplication, with up to 17 relevant text chunks selected for the final context. We observed a warning indicating that rerank is enabled but no reranker is configured; this is slated for resolution in the Action Plan. With response caching enabled, repeated queries complete significantly faster (sub-second wall time versus several seconds when cold). Vector similarity over 256 entity embeddings (1024-d) completes in tens of milliseconds on a laptop.

## 5 Challenges & Architecture Evolution

### 5.1 Limitations of Pure LightRAG

While LightRAG excels at semantic understanding and complex reasoning, we identified three critical limitations for production deployment:

**1. Inability to Leverage User Preferences** The system cannot effectively utilize structured user information stored in the `itineraries` table (budget, family composition, preferred categories, activity intensity). Vector-based semantic retrieval dilutes hard constraints—e.g., a query for “\$30 budget restaurants” may retrieve \$100 upscale venues due to semantic similarity in the “Tampa restaurants” context.

**2. Loss of Structured Query Precision** Converting structured database fields (e.g., `rating = 4.5`, `budget_cents ≤ 5000`) into natural language text and then embedding them sacrifices the precision of SQL-style filtering. A traditional indexed query `WHERE rating ≥ 4.0 AND budget_cents ≤ 5000` completes in <10 ms with guaranteed constraint satisfaction, whereas LightRAG’s vector retrieval (50–200 ms) yields approximate matches without guarantees.

### 5.2 Proposed Hybrid Architecture

To address these challenges, we propose a **hybrid PostgreSQL + LightRAG architecture** that combines the strengths of structured database queries with semantic graph-based reasoning:

## Design Principles

- **PostgreSQL:** Handles hard constraints (city, budget, rating, family-friendliness) via indexed queries.
- **LightRAG:** Handles semantic understanding (“romantic atmosphere,” “suitable for couples”) and generates personalized recommendation narratives.

## Three-Stage Retrieval Pipeline

### 1. Stage 1: PostgreSQL Pre-filtering

- Apply hard constraints extracted from user preferences:

```
SELECT p.* FROM places p
JOIN place_categories pc ON p.id = pc.place_id
WHERE p.city = 'Tampa'
    AND p.budget_cents <= 5000
    AND p.rating >= 4.0
    AND p.is_family_friendly = TRUE
ORDER BY p.popularity_score DESC
LIMIT 200;
```

- Returns 50–200 candidate POIs guaranteed to satisfy all hard constraints.

### 2. Stage 2: LightRAG Semantic Retrieval

- Load a pre-built city-specific knowledge graph (see below).
- Filter the KG to include only candidate POIs from Stage 1.
- Perform vector-based semantic search within this subgraph to identify the top-10 POIs that best match soft constraints (e.g., “romantic atmosphere”).

### 3. Stage 3: LLM Generation

- Assemble context from the top-10 semantic matches.
- Generate personalized recommendations with justifications via the LLM.

**Pre-built City Knowledge Graphs** To eliminate real-time KG construction overhead, we adopt an **offline pre-build strategy**:

- **Offline Phase** (daily/weekly batch): For each city (Tampa, Miami, Orlando, etc.), fetch all POIs from Google Places API, generate enhanced natural-language descriptions, and build a city-specific KG using LightRAG. Store the resulting graph and vector indices on disk.
- **Online Phase** (user query): Load the relevant city KG from cache (typically <10 ms if in memory) and proceed directly to Stage 2 filtering and retrieval.

## Performance Comparison(Anticipated)

Metric	Pure LightRAG	Hybrid Arch.	Improvement
Query Latency	12–25 s	2.5–5.5 s	78%
Hard Constraint Satisfaction	~60%	100%	+40 pp
Google Map API Cost	Per query	Daily/weekly batch	–95%
Data Freshness	Real-time	1-day delay	Acceptable

The hybrid approach delivers a 78% latency reduction while ensuring 100% constraint compliance and dramatically lowering API costs. The 1-day data staleness is acceptable for travel recommendations, as POI attributes (hours, ratings) change infrequently.

## 6 Action Plan

### 6.1 Phase 1: Validation & Prototyping (2 weeks)

1. **Database Schema Migration:** Execute schema updates to add structured fields (city, rating, budget\_cents, is\_family\_friendly, etc.) and create composite indexes for fast filtering.
2. **Enhanced Data Collection:** Implement multi-layer natural-language description generation for Google Places data (explicitly stating budget range, target audience, amenities) to improve LLM entity extraction quality.
3. **Test Data Ingestion:** Collect and validate POI data for 3 pilot cities (Tampa, Miami, Orlando) covering ~500 POIs each.

### 6.2 Phase 2: Hybrid Retrieval Implementation (2 weeks)

1. **Stage 1 PostgreSQL Filter:** Develop the pre-filtering module that constructs dynamic SQL queries from user preference objects (itineraries table).
2. **Stage 2 Semantic Retrieval:** Implement KG sub-graph extraction logic to filter pre-built city graphs by candidate POI IDs and perform local-mode retrieval.
3. **Stage 3 LLM Generation:** Integrate user preference context into the LLM prompt for personalized narrative generation.
4. **Benchmarking:** Compare pure LightRAG vs. hybrid pipeline on query latency, constraint satisfaction rate, and recommendation quality.

### 6.3 Phase 3: User Preference Integration (1-2 weeks)

1. **Preference Extraction:** Query the itineraries and itinerary\_preferred\_categories tables to extract user preferences (budget, family composition, activity intensity, etc.).
2. **Query Contextualization:** Automatically expand user queries with preference-based constraints (e.g., “Recommend restaurants in Tampa” → “...within \$50 budget, suitable for families with children, preferably Italian/Seafood”).
3. **Validation & Post-Processing:** Implement constraint verification to ensure LLM-generated recommendations comply with hard limits; trigger re-retrieval if violations occur.

### 6.4 Phase 4: Production Deployment (2-3 weeks)

1. **RESTful API:** Deploy a FastAPI-based service exposing /recommend endpoints that accept city, query, and itinerary\_id parameters.
2. **KG Update Mechanism:** Implement incremental KG updates via scheduled batch jobs (nightly or weekly) to refresh city graphs with new or updated POIs.
3. **Caching & Optimization:** Use Redis for in-memory KG caching, configure database connection pooling, and enable asynchronous processing for concurrent query handling.
4. **Testing & Documentation:** Conduct unit and integration tests; generate Swagger-based API documentation and deployment runbooks.

### 6.5 Future Extensions

- **Reranking Integration:** Wire lightweight rerankers (e.g., Cohere Rerank) to improve entity/chunk selection quality post-retrieval.

- **Multi-City Expansion:** Scale the offline KG-building pipeline to cover all major Florida cities and eventually extend to other states.
- **Real-Time Feedback Loop:** Incorporate user interaction signals (clicks, bookings, ratings) to fine-tune retrieval parameters and LLM prompts.
- **Advanced Personalization:** Leverage collaborative filtering or neural recommendation models to augment graph-based retrieval with user behavior patterns.

## References

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Vladimir Karpukhin, Naman Goyal, Ankush Kulkarni, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Luke Zettlemoyer. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

Yiming Zhu, Yangyi Chen, Zixuan Ma, Zhijing Jin, Wayne Xin Zhao, and Ji-Rong Wen. Lightrag: Simple and fast retrieval-augmented generation, 2024.